

Mini-C Challenges Submission

Hello, here is where you will see my C submissions!

Part 1

Print "Hello, NAME" where NAME is input from the keyboard.

```
In [ ]: #include <stdio.h>

int main() {
    char fullName[30];

    printf("Type your full name: \n");
    fgets(fullName, sizeof(fullName), stdin);

    printf("Hello %s", fullName);

    return 0;
}
// I used fgets because it says to use keyboard input in the instructions an
```

```
$ gcc part1.c -o part1 && ./part1
Type your full name:
Trevor Hurst
Hello Trevor Hurst
```

Part 2

Implement Archimedes algorithm to estimate pi for inscribed/circumscribed polygons with n sides, up to 100, doubling n at each step, and time it.

I used `$ sudo perf stat ./part2` for timing in this case. Yes, you can use a `clock_t` variable and `time.h` but, I didn't want the timing to mess with the code execution, `perf stat` also gives us more performance statistics

```
In [ ]: #include <stdio.h>
#include <math.h>

int sides = 3;
float len_inner = 0.8660254; // 0.5*sqrt(3);
```

```

float len_outer = 1.7320508; // sqrt(3);
float perimeter_inner;
float perimeter_outer;

int main() {
    while (sides < 100) {
        perimeter_outer = len_outer*sides;
        perimeter_inner = len_inner*sides;

        len_outer = len_outer / (1 + sqrt(1 + len_outer * len_outer));
        len_inner = sqrt(0.5 - 0.5 * sqrt(1 - len_inner * len_inner));

        sides *= 2;
    }
    printf("pi bounds: %f <= pi <= %f\n", perimeter_inner, perimeter_outer);
}

```

```

$ sudo perf stat ./part2
pi bounds: 3.141038 <= pi <= 3.142715

```

Performance counter stats for './part2':

	0.33 msec task-clock	#	0.334
CPUs utilized			
/sec	0 context-switches	#	0.000
/sec	0 cpu-migrations	#	0.000
K/sec	70 page-faults	#	213.389
GHz	1,212,428 cycles	#	3.696
frontend cycles idle	26,629 stalled-cycles-frontend	#	2.20%
backend cycles idle	232,277 stalled-cycles-backend	#	19.16%
insn per cycle	775,008 instructions	#	0.64
stalled cycles per insn		#	0.30
M/sec	176,379 branches	#	537.675
<not counted>	branch-misses		
(0.00%)			
0.000983511 seconds time elapsed			
0.001228000 seconds user			
0.000000000 seconds sys			

Part 3

Implement matrix – vector multiplication. Read in the following text file (mv.txt) which contains the matrix and vector to be multiplied. Print your answer to the screen and time the computation. The format of mv.txt is: line 1 contains numrows, numcols. The next numrows contains the rows of the matrix, each with numcols integer values. The next line contains the length of the vector. The next line contains the vector of that length.

In [3]:

```
#include <stdio.h>
#include <string.h>
#include <math.h>

int main() {

    FILE *fptr;
    int numrows, numcols;
    fptr = fopen("mv.txt", "r");

    fscanf(fptr, "%d %d", &numrows, &numcols);

    int matrix[numrows][numcols];
    for (int i = 0; i < numrows; i++) {
        for(int j = 0; j < numcols; j++) {
            fscanf(fptr, "%d", &matrix[i][j]);
        }
    }

    int vectorlen;
    fscanf(fptr, "%d", &vectorlen);
    int vector[vectorlen];

    for (int i = 0; i < vectorlen; i++) {
        fscanf(fptr, "%d", &vector[i]);
    }

    int result[numrows];
    memset(result, 0, sizeof(result));

    for (int i = 0; i < numrows; i++) {
        for (int j = 0; j < numcols; j++) {
            result[i] += vector[j]*matrix[i][j];
        }
    }

    for (int i = 0; i < numrows; i++) {
        printf("%d ", result[i]);
    }
    fclose(fptr);

    return 0;
}
```

```
18817 18431 5967 16686 25429 18817 16835 17175 25000 14060
```

```
$ sudo perf stat ./part3
18817 18431 5967 16686 25429 18817 16835 17175 25000 14060
Performance counter stats for './part3':
```

	#	Value
CPUs utilized	0.30	msec task-clock
/sec	0	context-switches
/sec	0	cpu-migrations
K/sec	60	page-faults
GHz	1,106,680	cycles
frontend cycles idle	23,867	stalled-cycles-frontend
backend cycles idle	209,091	stalled-cycles-backend
insn per cycle	740,057	instructions
stalled cycles per insn	165,650	branches
M/sec	<not counted>	branch-misses
(0.00%)		
0.001007680 seconds time elapsed		
0.001234000 seconds user		
0.000000000 seconds sys		

Part 4

Compare the speed of *, /, sqrt, sin operations/functions.

I am going to make multiple programs for this part, that way I can get the nice cool performance statistics for my code!

I decided to go with multiplying and dividing by 3 and making sure to turn off compiler optimizations so we're comparing apples-to-apples.

I had been doing $i*i$ and i/i but, the compiler was able to see that a lot of the operations could be changed to bit shifting which is way faster!

```
In [11]: // multiplication.c
#include <math.h>
```

```
int main() {
    int j = 0;
    for (int i = 0; i < 1000000000; i++) {
        j = i*3;
    }
}
```

\$ sudo perf stat ./multiplication

Performance counter stats for './multiplication':

CPUs utilized /sec GHz frontend cycles idle backend cycles idle insn per cycle stalled cycles per insn G/sec of all branches	693.11 msec task-clock # 0.999 1 context-switches # 1.443 0 cpu-migrations # 0.000 50 page-faults # 72.139 2,595,447,558 cycles # 3.745 (83.27%) 28,501 stalled-cycles-frontend # 0.00% (83.26%) 1,837,508 stalled-cycles-backend # 0.07% (83.26%) 8,001,252,232 instructions # 3.08 # 0.00 branches # 1.443 (83.48%) 35,998 branch-misses # 0.00% (83.46%)
0.694004320 seconds time elapsed	
0.694155000 seconds user	
0.000000000 seconds sys	

In []: // division.c

```
#include <math.h>
int main() {
    int j = 0;
    for (int i = 0; i < 1000000000; i++) {
        j = i/3;
    }
}
```

\$ sudo perf stat ./division

Performance counter stats for './division':

```

          802.93 msec task-clock      # 0.999
CPUs utilized
          1      context-switches    # 1.245
/sec
          0      cpu-migrations     # 0.000
/sec
          51     page-faults        # 63.518
/sec
          3,006,455,583    cycles      # 3.744
GHz
          38,173    (83.06%)
frontend cycles idle      stalled-cycles-frontend # 0.00%
          (83.06%)
          2,806,425    stalled-cycles-backend # 0.09%
backend cycles idle      (83.53%)
          12,000,690,795   instructions # 3.99
insn per cycle
                           # 0.00
stalled cycles per insn  (83.56%)
          1,000,527,651    branches    # 1.246
G/sec
                           (83.56%)
          49,178     branch-misses # 0.00%
of all branches          (83.23%)

0.803838726 seconds time elapsed

0.804003000 seconds user
0.000000000 seconds sys

```

```
In [ ]: // sqrt.c
#include <math.h>
int main(){
    int j = 0;
    for (int i = 0; i < 1000000000; i++) {
        j = sqrt(i);
    }
}
```

```

$ sudo perf stat ./sqrt

Performance counter stats for './sqrt':

          2,675.32 msec task-clock      # 1.000
CPUs utilized
          4      context-switches    # 1.495
/sec
          0      cpu-migrations     # 0.000
/sec
          61     page-faults        # 22.801
/sec
          10,017,423,521    cycles      # 3.744

```

```

GHz          (83.26%)
    75,012    stalled-cycles-frontend # 0.00%
frontend cycles idle      (83.25%)
    7,201,751  stalled-cycles-backend # 0.07%
backend cycles idle      (83.38%)
    20,004,402,439  instructions      # 2.00
insn per cycle
                                         # 0.00
stalled cycles per insn (83.40%)
    6,001,188,618  branches      # 2.243
G/sec          (83.40%)
    146,444     branch-misses # 0.00%
of all branches      (83.31%)

```

2.676314956 seconds time elapsed

2.676377000 seconds user
0.000000000 seconds sys

```
In [ ]: // sin.c
#include <math.h>
int main(){
    int j = 0;
    for (int i = 0; i < 1000000; i++){
        j = sin(i);
    }
}
```

\$ sudo perf stat ./sin

Performance counter stats for './sin':

```

            36,694.59 msec task-clock      # 1.000
CPUs utilized
    /sec          81    context-switches      # 2.207
                  0    cpu-migrations      # 0.000
    /sec          65    page-faults      # 1.771
    /sec          137,404,519,046  cycles      # 3.745
GHz          (83.33%)
    1,729,605    stalled-cycles-frontend # 0.00%
frontend cycles idle      (83.33%)
    91,795,142    stalled-cycles-backend # 0.07%
backend cycles idle      (83.33%)
    336,083,010,078  instructions      # 2.45
insn per cycle
                                         # 0.00
stalled cycles per insn (83.33%)
    23,884,801,515  branches      # 650.908

```

```
M/sec          (83.33%)
      12,191,316   branch-misses    #
of all branches      (83.33%)

36.697755393 seconds time elapsed

36.695391000 seconds user
0.000000000 seconds sys
```

Part 5

Use the attached code snippets as a basis for comparing the performance of row-major vs. column major computations. One snippet uses a static allocation for the array, the other allocates the array dynamically. Do a little experimentation with each approach. Vary the size of the square array from 128 X 128 on up, doubling it in size each time. Chart your results. Is there a difference in performance or behavior between static and dynamic? Between row-major and column-major? In terms of the latter, valgrind has been installed on your Linux VM, and its cachegrind tool facility may help provide some insights. Do some OSINT research to learn about valgrind....

```
In [12]: # include <time.h>
# include <math.h>
# include <stdio.h>
# include <stdlib.h>

int main(int argc, char **argv) {
    int i,j;
    int n = 128;
    double sum;
    clock_t end, start;
    double arr[128][128];

    // THIS FILLS THE MATRIX WITH NUMBERS
    for (i=0; i<n; i++){
        for (j=0; j<n; j++){
            arr[i][j] = (double) rand() / RAND_MAX;
        }
    }

    sum = 0;

    start = clock();

    // ROW MAJOR WORK
    // YOU'LL NEED TO TIME IT
    for (i = 0; i<n; i++){ // iterate over rows
        for (j = 0; j<n; j++){ // iterate over columns
            sum += arr[i][j];
        }
    }
```

```

        end = clock();

// NOTE: YOU'LL NEED TO PROVIDE MEANING TO end AND start
printf("Row Major: sum = %lf and Clock Ticks are %ld\n", sum, end-start);

//ADD YOUR COLUMN MAJOR WORK
// YOU'LL NEED TO TIME IT

    return 0;
}

```

Row Major: sum = 8143.619987 and Clock Ticks are 55

```

$ valgrind --tool=cachegrind ./loopsstaticshell
==2497347== Cachegrind, a cache and branch-prediction
profiler
==2497347== Copyright (C) 2002-2017, and GNU GPL'd, by
Nicholas Nethercote et al.
==2497347== Using Valgrind-3.18.1 and LibVEX; rerun with -h
for copyright info
==2497347== Command: ./loopsstaticshell
==2497347==
--2497347-- warning: L3 cache found, using its data for the
LL simulation.

Row Major: sum = 8143.619987 and Clock Ticks are 877
Column Major: sum = 16287.239974 and Clock Ticks are 808
==2497347==
==2497347== I refs: 1,862,479
==2497347== I1 misses: 1,484
==2497347== LLi misses: 1,457
==2497347== I1 miss rate: 0.08%
==2497347== LLi miss rate: 0.08%
==2497347==
==2497347== D refs: 840,028 (661,300 rd + 178,728
wr)
==2497347== D1 misses: 22,919 ( 20,193 rd + 2,726
wr)
==2497347== LLd misses: 3,906 ( 1,376 rd + 2,530
wr)
==2497347== D1 miss rate: 2.7% ( 3.1% +
1.5% )
==2497347== LLd miss rate: 0.5% ( 0.2% +
1.4% )
==2497347==
==2497347== LL refs: 24,403 ( 21,677 rd + 2,726
wr)
==2497347== LL misses: 5,363 ( 2,833 rd + 2,530
wr)
==2497347== LL miss rate: 0.2% ( 0.1% +
1.4% )

```

```
In [14]: # include <time.h>
# include <math.h>
# include <stdio.h>
# include <stdlib.h>

int main(int argc, char **argv) {
    int i,j;
    int n = 128;
    double sum;
    clock_t end, start;
    double *arr = malloc(n*n*sizeof(double));

    // THIS FILLS THE MATRIX WITH NUMBERS
    for (i=0; i<n; i++){
        for (j=0; j<n; j++){
            arr[i*n+j] = (double) rand() / RAND_MAX;
        }
    }

    sum = 0;

    start = clock();
    // ROW MAJOR WORK
    // YOU'LL NEED TO TIME IT
    for (i = 0; i<n; i++){ // iterate over rows
        for (j = 0; j<n; j++){ // iterate over columns
            sum += arr[i*n + j];
        }
    }
    end = clock();

    printf("Row Major: sum = %lf and Clock Ticks are %ld\n", sum, end-start);

    //ADD YOUR COLUMN MAJOR WORK
    // YOU'LL NEED TO TIME IT

    return 0;
}
```

Row Major: sum = 8143.619987 and Clock Ticks are 54

```
$ valgrind --tool=cachegrind ./loopsdynamicshell
==2497352== Cachegrind, a cache and branch-prediction
profiler
==2497352== Copyright (C) 2002-2017, and GNU GPL'd, by
Nicholas Nethercote et al.
==2497352== Using Valgrind-3.18.1 and LibVEX; rerun with -h
for copyright info
==2497352== Command: ./loopsdynamicshell
==2497352==
--2497352-- warning: L3 cache found, using its data for the
LL simulation.
Row Major: sum = 8143.619987 and Clock Ticks are 976
Column Major: sum = 16287.239974 and Clock Ticks are 903
```

```

==2497352==
==2497352== I refs: 2,009,802
==2497352== I1 misses: 1,483
==2497352== LLi misses: 1,453
==2497352== I1 miss rate: 0.07%
==2497352== LLi miss rate: 0.07%
==2497352==
==2497352== D refs: 938,290 (759,538 rd + 178,752
wr)
==2497352== D1 misses: 22,889 ( 20,161 rd + 2,728
wr)
==2497352== LLd misses: 3,975 ( 1,338 rd + 2,637
wr)
==2497352== D1 miss rate: 2.4% ( 2.7% +
1.5% )
==2497352== LLd miss rate: 0.4% ( 0.2% +
1.5% )
==2497352==
==2497352== LL refs: 24,372 ( 21,644 rd + 2,728
wr)
==2497352== LL misses: 5,428 ( 2,791 rd + 2,637
wr)
==2497352== LL miss rate: 0.2% ( 0.1% +
1.5% )

```

Running with multiple and graphing

 Here we can clearly see there is a large discrepancy in row major vs column major addition!

Part 6

Write a program that accepts a string input from stdio and sends it to a function that transforms it according a transposition function passed in to it as an argument. The function will print out the string, transform it, and then print out the result. The transposition function, you can assume, simply shuffles the existing characters in the string. Build a transposition function that reverses the string and apply it. Where appropriate and possible, use dynamic allocation and pointer arithmetic to get the job done.

```
In [ ]: #include <stdio.h>
#include <stdlib.h>
#include <string.h>

// This makes the code more reusable, we can have a bunch of functions and s
typedef void (*TranspositionFunc) (char *);
```

```

void reverse_string(char *str) {
    if (str == NULL) return;

    char *start = str;
    char *end = str + strlen(str) - 1;

    char temp;

    while (start < end) {
        temp = *start;
        *start = *end;
        *end = temp;

        start++;
        end--;
    }
}

// Accepts the data (str) and the function to apply to it (func)
void apply_and_print(char *str, TranspositionFunc func) {
    printf("\n--- Transposition Report ---\n");
    printf("Original String: %s\n", str);

    func(str);

    printf("Transformed String: %s\n", str);
    printf("-----\n");
}

int main() {
    size_t buffer_size = 1024;

    // Dynamic Allocation: Allocate memory on the heap for the input
    char *input_buffer = (char *)malloc(buffer_size * sizeof(char));

    if (input_buffer == NULL) {
        fprintf(stderr, "Memory allocation failed\n");
        return 1;
    }

    printf("Enter a string to transpose: ");

    // Read input from stdin
    if (getline(&input_buffer, &buffer_size, stdin) != -1) {
        // Strip the newline character
        char *newline = strchr(input_buffer, '\n');
        if (newline) *newline = '\0';

        // Pass in the string and the function reverse_string
        apply_and_print(input_buffer, reverse_string);
    } else {
        printf("Error reading input.\n");
    }

    // Clean up heap memory
    free(input_buffer);
}

```

```
    return 0;
}
```

```
$ ./part6
Enter a string to transpose: Hello World! -Trevor

--- Transposition Report ---
Original String:    Hello World! -Trevor
Transformed String: roverT- !dlroW olleH
-----
```

Final Questions

1. For the Hello Name challenge:

- a. Did you try passing your name as an argument from the command line or did you use `scanf`? Why?

I used `fgets` to grab input from the terminal. This is how I learned to do it, but I can see how it is less parallelizable than grabbing from a file or using `clargs`!

- b. How did you manage or allocate the strings? (Static or dynamic)

I used static allocation, I wasn't quite as familiar with dynamic string allocation in C until midway through this assignment!

2. For Archimedes algorithm

- a. How did you time your program?

I timed it using `perf stat` I did it this way to keep overhead in my executable lower!

- b. Were there any issues with precision and/or convergence that you noticed?

There were issues if I ran the code past the suggested 100 limit but, precision up until that point was fine, there was probably some floating point arithmetic errors but, nothing that made my answers too far off!

3. For Matrix-vector multiplication

- a. How did you allocate and access your matrix?

I allocated it and accessed it dynamically by taking in file information.

- b. Were there any challenges in reading in the file?

It took me a minute to stop and think about the best way to read the file in, but fscanf was pretty elegant!

c. Was there anything special about the actual computation?

Not in particular, it took looking back at linear algebra notes and using an online matrix calculator to sanity check my results before I was comfortable submitting.

d. What was your strategy for timing?

Again to reduce overhead in my program, I utilized `perf stat` which was able to give very detailed and helpful timing information

4. For measuring the speed of arithmetic computations

a. What was your timing strategy?

Again to reduce overhead in my program, I utilized `perf stat` which was able to give very detailed and helpful timing information

b. Are all arithmetic operations created equal?

No not all arithmetic operations are equal! The compiler can try and do some optimizations like if it sees a div with a multiple of 2 it can change that to a left shift but, the operations without dedicated operators are much slower! Particularly sin! A way to combat the slowness of something like sin is to use pre-calculated tables or values!

5. For the row-major/column-major exercise

a. What did you observe about differences in program behavior in static vs dynamic allocation of arrays, and how do you explain it?

We had to be much more explicit in the static arrays, changing the array size and n, there were some small performance increases with the dynamic but that was just in the startup phase, once we got into the loop, the arrays were identical.

b. What did you observe about differences in program behavior in row-major vs. column major computations and how do you explain it?

Row-major order ended up taking over in the speed game, likely due to the cache's spatial locality! It is much easier to access memory close to what you just accessed than memory further away! We had much fewer cache misses with row major computations!